

approach

what it needs

what it returns

generate and select plausible patients

free: which candidates are kept

a range per observable



a set of patients,
none more common than another

reweight a fixed cloud

free: prevalence weights on the cloud

observed marginals



one weighting,
and no band around it

fit patients individually

free: one θ per real patient

individual patient data



one posterior *per patient*

this model

free: (μ, ω) and the discrepancy

cohort summaries, with n



a *distribution over populations*,
one per posterior draw